

Data manipulation with **dplyr**

Hadley Wickham

@hadleywickham

Chief Scientist, RStudio



June 2014

Data analysis is a process
by which data becomes
understanding, knowledge
and insight

Data analysis is the process
by which data becomes
understanding, knowledge
and insight

Tidy



Transform

Visualise

Surprises, but doesn't scale

Model

Scales, but doesn't (fundamentally) surprise

Tidy

tidyr



Transform

dplyr

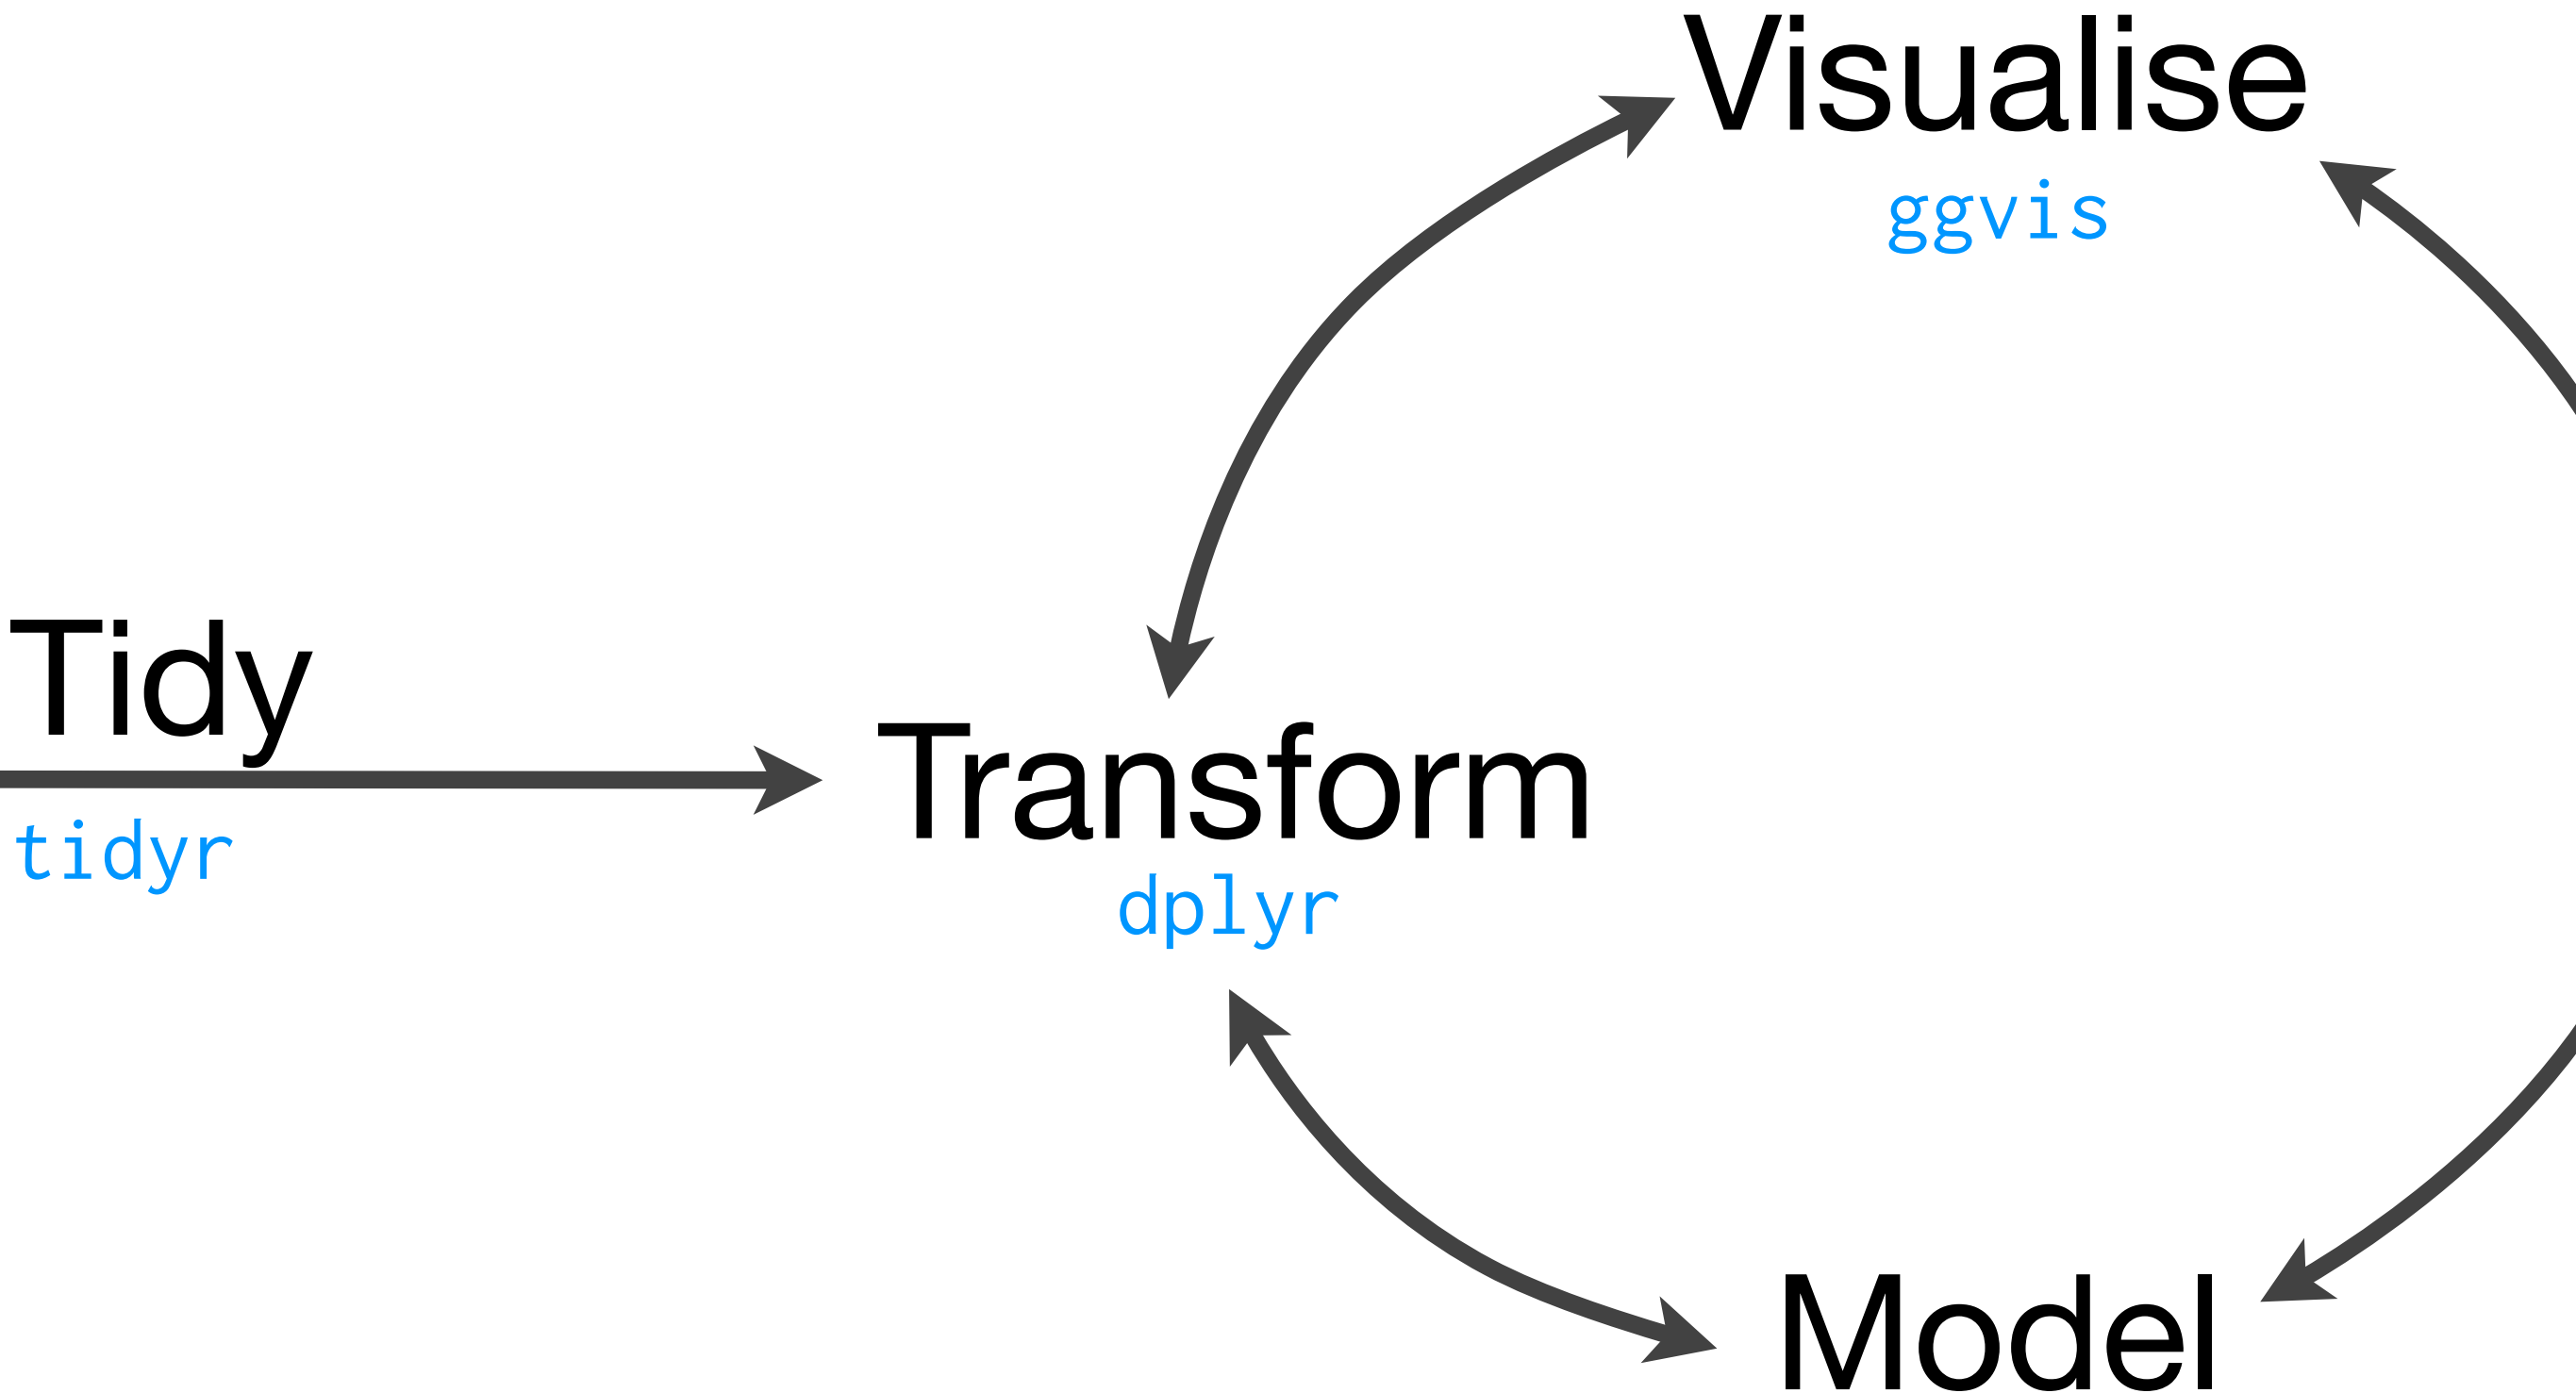


Visualise

ggvis



Model



1. Flights data
2. One table verbs & grouped summaries
3. Data pipelines
4. Grouped mutate/filter & window functions
5. Joins (two table verbs)
6. Do
7. Databases

The bad news:
It's going to be
frustrating





3:17 AM

The good news:
Frustration is
typical and
temporary

Flights data

Rstudio projects

- Isolate code and results from different projects. Restart where you left off.
- Double-click `dp1yr-tutorial.Rproj` file to open. (One R file for each section)
- (If you don't use RStudio, just change working directories)

Flights data

- `flights` [227,496 x 14]. Every flight departing Houston in 2011.
- `weather` [8,723 x 14]. Hourly weather data.
- `planes` [2,853 x 9]. Plane metadata.
- `airports` [3,376 x 7]. Airport metadata.

```
library(dplyr)
library(ggplot2)

flights <- tbl_df(read.csv("flights.csv",
  stringsAsFactors = FALSE))
flights$date <- as.Date(flights$date)

weather <- tbl_df(read.csv("weather.csv",
  stringsAsFactors = FALSE))
weather$date <- as.Date(weather$date)

planes <- tbl_df(read.csv("planes.csv",
  stringsAsFactors = FALSE))

airports <- tbl_df(read.csv("airports.csv",
  stringsAsFactors = FALSE))
```

Your turn

Introduce yourself to your neighbour.

What questions might you want to answer with this data?

**One table
verbs**

- **filter:** keep rows matching criteria
- **select:** pick columns by name
- **arrange:** reorder rows
- **mutate:** add new variables
- **summarise:** reduce variables to values

Structure

- First argument is a data frame
- Subsequent arguments say what to do with data frame
- Always return a data frame
- (Never modify in place)


```
df <- data.frame(  
  color = c("blue", "black", "blue", "blue", "black"),  
  value = 1:5)
```

df

color	value
blue	1
black	2
blue	3
blue	4
black	5



color	value
blue	1
blue	3
blue	4

```
filter(df, color == "blue")
```

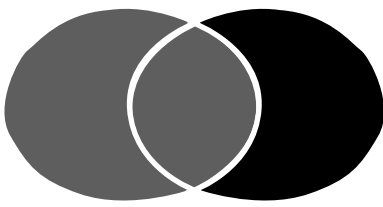
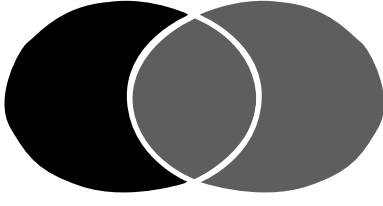
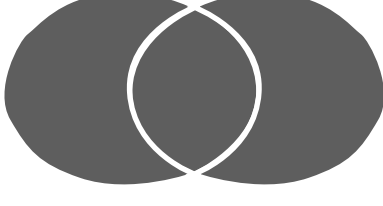
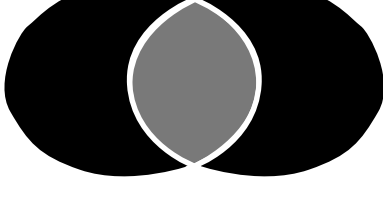
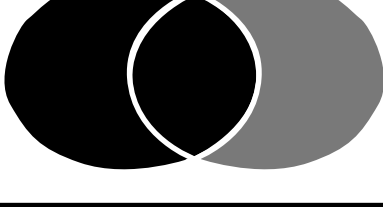
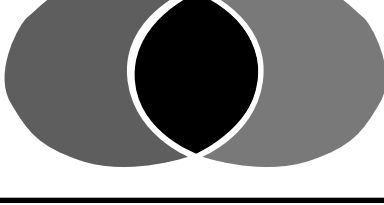
df

color	value
blue	1
black	2
blue	3
blue	4
black	5



color	value
blue	1
blue	4

```
filter(df, value %in% c(1, 4))
```

	a
	b
	a b
	a & b
	a & !b
	xor(a, b)

x > 1

x >= 1

x < 1

x <= 1

x != 1

x == 1

x %in% ("a", "b")

Find all flights:

To SFO or OAK

In January

Delayed by more than an hour

That departed between midnight and five am.

Where the arrival delay was more than twice the departure delay

```
filter(flights, dest %in% c("SFO", "OAK"))  
filter(flights, dest == "SFO" | dest == "OAK")  
# Not this!  
filter(flights, dest == "SFO" | "OAK")  
  
filter(flights, date < "2001-02-01")  
  
filter(flights, hour >= 0, hour <= 5)  
filter(flights, hour >= 0 & hour <= 5)  
  
filter(flights, dep_delay > 60)  
  
filter(flights, arr_delay > 2 * dep_delay)
```

df

color	value
blue	1
black	2
blue	3
blue	4
black	5



color
blue
black
blue
blue
black

```
select(df, color)
```

df

color	value
blue	1
black	2
blue	3
blue	4
black	5



value
1
2
3
4
5

```
select(df, -color)
```


Your turn

Read the help for `select()`. What other ways can you select variables?

Write down three ways to select the two delay variables.

```
select(flights, arr_delay, dep_delay)
select(flights, arr_delay:dep_delay)
select(flights, ends_with("delay"))
select(flights, contains("delay"))
```

df

color	value
4	1
1	2
5	3
3	4
2	5



color	value
1	2
2	5
3	4
4	1
5	3

```
arrange(df, color)
```

df

color	value
4	1
1	2
5	3
3	4
2	5



color	value
5	3
4	1
3	4
2	5
1	2

```
arrange(df, desc(color))
```

Your turn

Order the flights by departure date and time.

Which flights were most delayed?

Which flights caught up the most time during the flight?

```
arrange(flights, date, hour, minute)
```

```
arrange(flights, desc(dep_delay))
```

```
arrange(flights, desc(arr_delay))
```

```
arrange(flights, desc(dep_delay - arr_delay))
```

df

color	value
blue	1
black	2
blue	3
blue	4
black	5



color	value	double
blue	1	2
black	2	4
blue	3	6
blue	4	8
black	5	10

```
mutate(df, double = 2 * value)
```

df

color	value
blue	1
black	2
blue	3
blue	4
black	5



color	value	double	quadruple
blue	1	2	4
black	2	4	8
blue	3	6	12
blue	4	8	16
black	5	10	20

```
mutate(df, double = 2 * value,  
       quadruple = 2 * double)
```


Your turn

Compute speed in mph from time (in minutes) and distance (in miles). Which flight flew the fastest?

Add a new variable that shows how much time was made up or lost in flight.

How did I compute hour and minute from dep?

(Hint: you may need to use `select()` or `View()` to see your new variable)

```
flights <- mutate(flights,  
  speed = dist / (time / 60))  
arrange(flights, desc(speed))  
  
mutate(flights, delta = dep_delay - arr_delay)  
  
mutate(flights,  
  hour = dep %/% 100,  
  minute = dep %% 100)
```

**Grouped
summarise**

df

color	value
blue	1
black	2
blue	3
blue	4
black	5



total
15

```
summarise(df, total = sum(value))
```

df

color	value
blue	1
black	2
blue	3
blue	4
black	5



color	total
blue	8
black	7

```
by_color <- group_by(df, color)
summarise(by_color, total = sum(value))
```

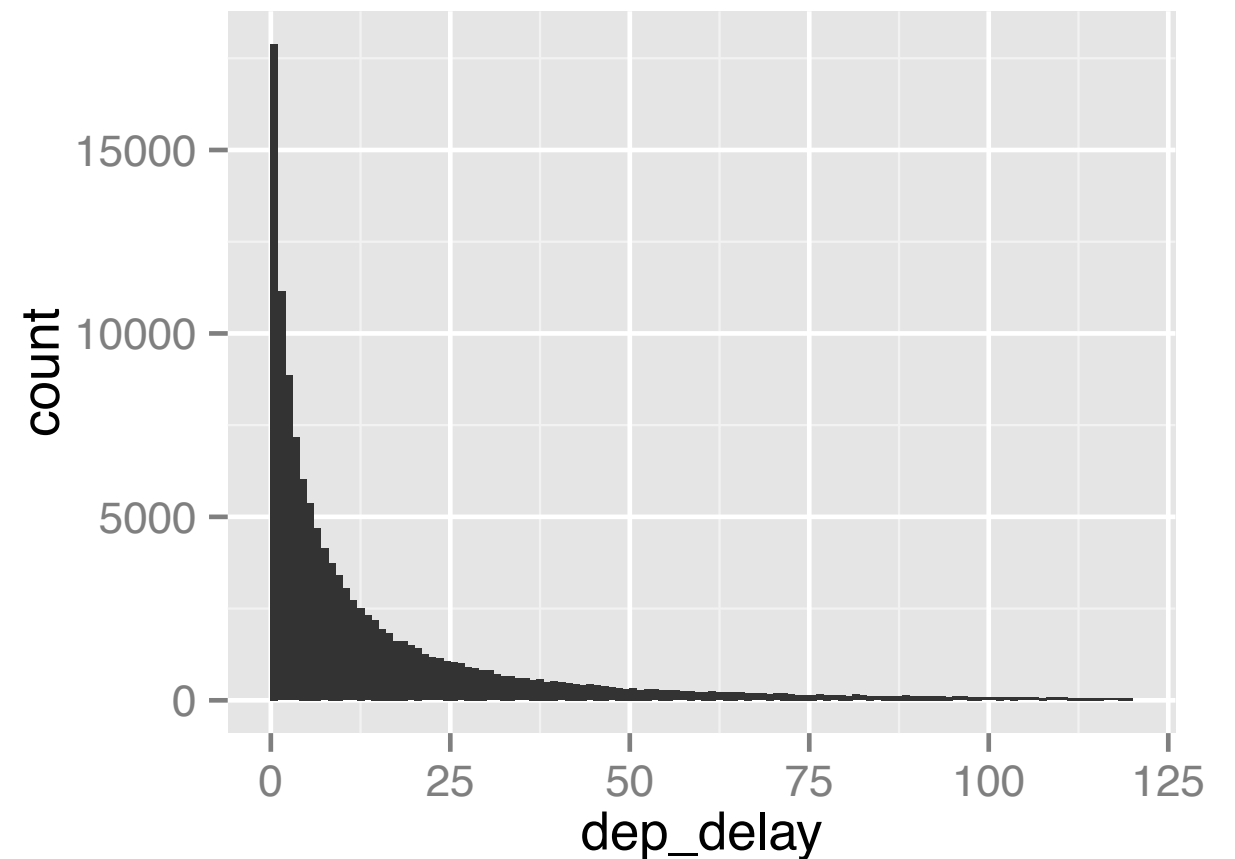
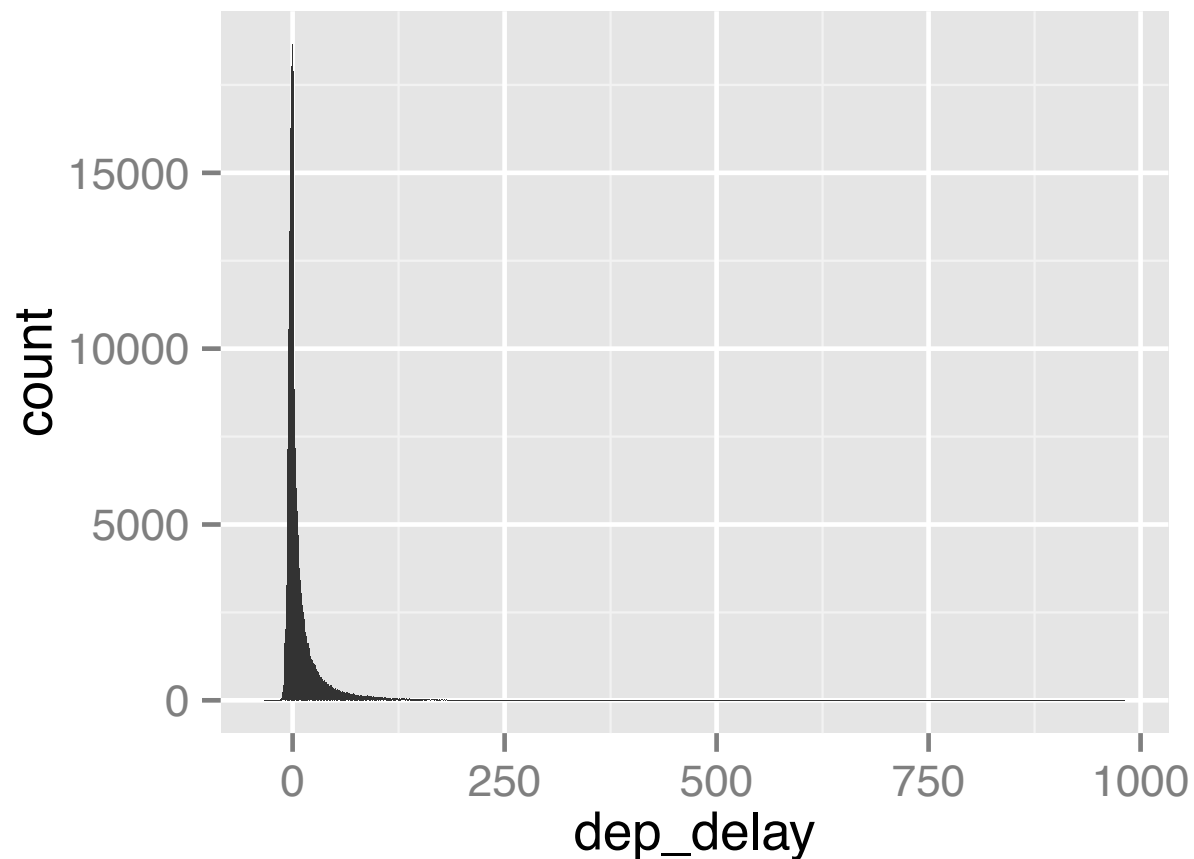
```
by_date <- group_by(flights, date)
by_hour <- group_by(flights, date, hour)
by_plane <- group_by(flights, plane)
by_dest <- group_by(flights, dest)
```

Summary functions

- `min(x)`, `median(x)`, `max(x)`,
`quantile(x, p)`
- `n()`, `n_distinct()`, `sum(x)`, `mean(x)`
- `sum(x > 10)`, `mean(x > 10)`
- `sd(x)`, `var(x)`, `iqr(x)`, `mad(x)`

Your turn

How might you summarise dep_delay for each day? Brainstorm for 2 minutes.




```
by_date <- group_by(flights, date)
delays <- summarise(by_date,
  mean = mean(dep_delay),
  median = median(dep_delay),
  q75 = quantile(dep_delay, 0.75),
  over_15 = mean(dep_delay > 15),
  over_30 = mean(dep_delay > 30),
  over_60 = mean(dep_delay > 60)
)
```

```
by_date <- group_by(flights, date)
delays <- summarise(by_date,
  mean = mean(dep_delay, na.rm = TRUE),
  median = median(dep_delay, na.rm = TRUE),
  q75 = quantile(dep_delay, 0.75, na.rm = TRUE),
  over_15 = mean(dep_delay > 15, na.rm = TRUE),
  over_30 = mean(dep_delay > 30, na.rm = TRUE),
  over_60 = mean(dep_delay > 60, na.rm = TRUE)
)
```

OR

```
by_date <- group_by(flights, date)
no_missing <- filter(flights, !is.na(dep))
delays <- summarise(no_missing,
  mean = mean(dep_delay),
  median = median(dep_delay),
  q75 = quantile(dep_delay, 0.75),
  over_15 = mean(dep_delay > 15),
  over_30 = mean(dep_delay > 30),
  over_60 = mean(dep_delay > 60)
)
```

Data pipelines

```
# Downside of functional interface is that it's
# hard to read multiple operations:
hourly_delay <- filter(
  summarise(
    group_by(
      filter(
        flights,
        !is.na(dep_delay)
      ),
      date, hour
    ),
    delay = mean(dep_delay),
    n = n()
  ),
  n > 10
)
```

```
# Solution: the pipe operator from magrittr
# x %>% f(y) -> f(x, y)

hourly_delay <- flights %>%
  filter(!is.na(dep_delay)) %>%
  group_by(date, hour) %>%
  summarise(delay = mean(dep_delay), n = n()) %>%
  filter(n > 10)

# Hint: pronounce %>% as then
```

Your turn

Create data pipelines to answer the following questions:

Which destinations have the highest average delays?

Which flights (i.e. carrier + flight) happen every day? Where do they fly to?

On average, how do delays (of non-cancelled flights) vary over the course of a day?
(Hint: $\text{hour} + \text{minute} / 60$)

```
flights %>%  
  group_by(dest) %>%  
  summarise(  
    arr_delay = mean(arr_delay, na.rm = TRUE),  
    n = n()) %>%  
  arrange(desc(arr_delay))  
  
# Nifty trick to see more data  
.Last.value %>% View()  
  
# It would be nice to plot these on a map...
```



```
flights %>%  
  group_by(carrier, flight, dest) %>%  
  tally(sort = TRUE) %>% # Save some typing  
  filter(n == 365)
```

```
flights %>%  
  group_by(carrier, flight, dest) %>%  
  summarise(n = n()) %>%  
  arrange(desc(n)) %>%  
  filter(n == 365)
```

Slightly different answer

```
flights %>%  
  group_by(carrier, flight) %>%  
  filter(n() == 365)
```

```
per_hour <- flights %>%
  filter(cancelled == 0) %>%
  mutate(time = hour + minute / 60) %>%
  group_by(time) %>%
  summarise(
    arr_delay = mean(arr_delay, na.rm = TRUE),
    n = n()
  )

qplot(time, arr_delay, data = per_hour)
qplot(time, arr_delay, data = per_hour, size = n) + scale_size_area()
qplot(time, arr_delay, data = filter(per_hour, n > 30), size = n) +
  scale_size_area()

ggplot(filter(per_hour, n > 30), aes(time, arr_delay)) +
  geom_vline(xintercept = 5:24, colour = "white", size = 2) +
  geom_point()
```

**Grouped
mutate / filter**

Groupwise variables

- Creating new variables within a group is also often useful.
- Sometime that's a combination of aggregation and recycling, e.g.
$$z = (x - \text{mean}(x)) / \text{sd}(x)$$
- Other times you need a **window function**
- More details in
`vignette("window-functions")`

Example:

```
planes <- flights %>%  
  filter(!is.na(arr_delay)) %>%  
  group_by(plane) %>%  
  filter(n() > 30)
```

```
planes %>%  
  mutate(z_delay =  
    (arr_delay - mean(arr_delay)) / sd(arr_delay)) %>%  
  filter(z_delay > 5)
```

```
planes %>% filter(min_rank(arr_delay) < 5)
```

Window functions

- Aggregation function:
 $n \text{ inputs} \rightarrow 1 \text{ output}$
- Window function:
 $n \text{ inputs} \rightarrow n \text{ outputs}$
- (Excludes functions that could operate row by row)

Types of window functions

- **Ranking and ordering**
- **Offsets: lead & lag**
- Cumulative aggregates
- Rolling aggregates

Your turn

What's the difference between `min_rank()`, `row_number()` and `dense_rank()`?

For each plane, find the two most delayed flights. Which of the three rank functions is most appropriate?


```
min_rank(c(1, 1, 2, 3))  
dense_rank(c(1, 1, 2, 3))  
row_number(c(1, 1, 2, 3))
```

```
flights %>% group_by(plane) %>%  
  filter(row_number(desc(arr_delay)) <= 2)
```

```
flights %>% group_by(plane) %>%  
  filter(min_rank(desc(arr_delay)) <= 2)
```

```
flights %>% group_by(plane) %>%  
  filter(dense_rank(desc(arr_delay)) <= 2)
```

```
daily <- flights %>%  
  group_by(date) %>%  
  summarise(delay = mean(dep_delay, na.rm = TRUE))  
  
# What's the day-to-day change?  
daily %>% mutate(delay - lag(delay))  
  
# If not ordered by date already  
daily %>% mutate(delay - lag(delay), order_by = date)
```

Other uses

- Was there a change? $x \neq \text{lag}(x)$
- Percent change? $(x - \text{lag}(x)) / x$
- Fold-change? $x / \text{lag}(x)$
- Previously false, now true? $!\text{lag}(x) \ \& \ x$

**Two table
verbs**

```
# Motivation: how can we show airport delays on  
# a map? Need to connect to airports dataset
```

```
location <- airports %>%  
  select(dest = iata, name = airport, lat, long)
```

```
flights %>%  
  group_by(dest) %>%  
  filter(!is.na(arr_delay)) %>%  
  summarise(  
    arr_delay = mean(arr_delay),  
    n = n()  
  ) %>%  
  arrange(desc(arr_delay)) %>%  
  left_join(location)
```

Joining datasets

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums
Stuart	bass
Pete	drums

+

name	band
John	T
Paul	T
George	T
Ringo	T
Brian	F

=

?

```
x <- data.frame(  
  name = c("John", "Paul", "George", "Ringo", "Stuart", "Pete"),  
  instrument = c("guitar", "bass", "guitar", "drums", "bass",  
    "drums")  
)  
  
y <- data.frame(  
  name = c("John", "Paul", "George", "Ringo", "Brian"),  
  band = c("TRUE", "TRUE", "TRUE", "TRUE", "FALSE")  
)
```

x

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums
Stuart	bass
Pete	drums

y

name	band
John	T
Paul	T
George	T
Ringo	T
Brian	F

+

=

name	instrument	band
John	guitar	T
Paul	bass	T
George	guitar	T
Ringo	drums	T

```
inner_join(x, y)
```


x

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums
Stuart	bass
Pete	drums

y

name	band
John	T
Paul	T
George	T
Ringo	T
Brian	F

+

=

name	instrument	band
John	guitar	T
Paul	bass	T
George	guitar	T
Ringo	drums	T
Stuart	bass	NA
Pete	drums	NA

```
left_join(x, y)
```

x

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums
Stuart	bass
Pete	drums

y

name	band
John	T
Paul	T
George	T
Ringo	T
Brian	F

+

=

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums

```
semi_join(x, y)
```

x

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums
Stuart	bass
Pete	drums

y

name	band
John	T
Paul	T
George	T
Ringo	T
Brian	F

+

=

name	instrument
Stuart	bass
Pete	drums

```
anti_join(x, y)
```

Type	Action
inner	Include only rows in both x and y
left	Include all of x, and matching rows of y
semi	Include rows of x that match y
anti	Include rows of x that don't match y

```
# Let's combine hourly delay data with weather  
# information
```

```
hourly_delay <- flights %>%  
  group_by(date, hour) %>%  
  filter(!is.na(dep_delay)) %>%  
  summarise(  
    delay = mean(dep_delay),  
    n = n()  
  ) %>%  
  filter(n > 10)  
delay_weather <- hourly_delay %>% left_join(weather)
```

Your turn

What weather conditions are associated with delays leaving in Houston?

Use graphics to explore.

```
qplot(temp, dep, data = delay_weather)
qplot(wind_speed, dep, data = delay_weather)
qplot(gust_speed, dep, data = delay_weather)
qplot(is.na(gust_speed), dep, data = delay_weather,
      geom = "boxplot")
qplot(conditions, dep, data = delay_weather,
      geom = "boxplot")
qplot(events, dep, data = delay_weather,
      geom = "boxplot")
```

Your turn

Are older planes more likely to be delayed? Explore the data and answer with a plot.

(Hint: I'd recommend by starting with some checking of the plane data)

Do

The workhorse function

- If one of the specialised verbs doesn't do what you need, you can use `do()`
- It's slower, but general purpose.
- Equivalent to `ddply()` and `dlply()`, and is particularly useful in conjunction with models

How it works

- Two variations: unnamed (for functions that return data frames), and named (for functions that return anything else)
- Uses a pronoun, ., to represent the current group

```
# Derived from http://stackoverflow.com/a/23341485/16632
library(dplyr)
library(zoo)
df <- data.frame(
  houseID = rep(1:10, each = 10),
  year = 1995:2004,
  price = ifelse(runif(10 * 10) > 0.50, NA, exp(rnorm(10 * 10)))
)

df %>%
  group_by(houseID) %>%
  do(na.locf(.))

df %>%
  group_by(houseID) %>%
  do(head(., 2))

df %>%
  group_by(houseID) %>%
  do(data.frame(year = . $year[1]))
```

```
# Named usage allows us to put any object into  
# a column: creates a "list-column". This is valid  
# in R, but data frame methods don't always expect.
```

```
df <- data.frame(x = 1:5)  
df$y <- list(1:2, 2:3, 3:4, 4:5, 5:6)
```

```
df  
str(df)
```

```
tbl_df(df)
```

```
# Doesn't work
```

```
df <- data.frame(  
  x = 1:5,  
  y = list(1:2, 2:3, 3:4, 4:5, 5:6)  
)
```

```
# Goal fit a linear model to each day, predicting  
# delay from time of day
```

```
usual <- flights %>%  
  mutate(time = hour + minute / 60) %>%  
  filter(hour >= 5, hour <= 20)
```

```
models <- usual %>%  
  group_by(date) %>%  
  do(  
    mod = lm(dep_delay ~ time, data = .)  
  )
```

```
# See 5-do.R for more details
```

Future work

- Labelling is still a little wonky
- Parallel? (like plyr)
- Better tools for working with models

Databases

Other data sources

- PostgreSQL, Greenplum, redshift
- MySQL, MariaDB
- SQLite
- MonetDB, BigQuery
- *Oracle, SQL Server, ImpalaDB*

Getting started

- Easiest to dip your toe in database waters with SQLite. No setup required!
- dplyr provides `copy_to()`, which makes it easy to get data from R into DB
- You can work with database tables just like data frames. dplyr translates the SQL for you.

```
hflights_db <- src_sqlite("hflights.sqlite3",  
  create = TRUE)
```

```
copy_to(  
  dest = hflights_db,  
  copy_to(hflights),  
  indexes = list(  
    c("date", "hour"),  
    "plane",  
    "dest",  
    "arr_delay"  
  ), temporary = FALSE  
)
```

Start with variables
needed to join tables

Default is to create
temporary tables

```
# DEMO
```

Learning SQL

- Learn how to use SELECT.
- Learn how indices work.
(<http://www.sqlite.org/queryplanner.html>)
- Learn how SELECT works.
(<http://tech.pro/tutorial/1555/10-easy-steps-to-a-complete-understanding-of-sql>)
- Make friends with an expert

When to use?

- Obviously, good idea to use if you data already in database. Better to pull from live db than to use static exports.
- If data fits in memory, using local data frame will always be faster. Only use DB for “big” data.
- Correct indexes are key to good filter + join performance. Talk to a DBA!

**Where
next**

```
browseVignettes(package = "dplyr")
```

```
# Translate plyr to dplyr
```

```
http://jimhester.github.io/plyrToDplyr/
```

```
# Common questions & answers
```

```
http://stackoverflow.com/questions/tagged/dplyr?  
sort=frequent
```